

File I/O

FILE

■ file

- unix에서의 파일은 단지 바이트들의 나열이다
- operating system은 파일에 어떤 포맷도 부과하지 않는다
- 파일의 내용은 바이트 단위로 주소를 줄 수 있다

■ file descriptor

- file descriptor는 0이나 양수이다
- file은 open이나 creat로 file descriptor가 새로 할당되고, read, write할 때 이를 이용한다
- 그러므로 user가 생성하는 첫 번째 파일의 file descriptor는 3이다.

■ file type

- regular file - 0 이상의 data block을 가진 일반적인 파일
- directory - 파일과 파일 이름을 mapping 시켜주는 파일
- special file - physical device를 file system에 mapping
- link - 파일을 다른 이름으로 연결
- symbolic link - 다른 파일을 가리키는 파일
- named pipe(FIFO) - process 간의 통신을 위한 파일

File Descriptor

❑ file descriptor

- all open files are referred to by file descriptors
- how to obtain file descriptor
 - return value of `open()`, `creat()`
- when we want to read or write a file, we identify the file with the file descriptor
- file descriptor is the index of user file descriptor table
- `STDIN_FILENO(0)`, `STDOUT_FILENO(1)`,
`STDERR_FILENO(2)` (`<unistd.h>`)

File Descriptor

❑ file descriptor (cont'd)

- standard input, output, error
 - STDIN_FILENO (== 0)
 - STDOUT_FILENO (== 1)
 - STDERR_FILENO (== 2)
- defined in <unistd.h>
- opened by the shell
 - not by the kernel

open()

■ synopsis

- `#include <sys/types.h>`
- `#include <sys/stat.h>`
- `#include <fcntl.h>`

- `int open(const char *pathname, int flags);`
- `int open(const char *pathname, int flags,
mode_t mode);`

open()

■ description

- attempts to open a file and return a file descriptor
- flags (defined in <fcntl.h>)
 - O_RDONLY, O_WRONLY or O_RDWR
 - access mode
 - O_CREAT
 - If the file does not exist it will be created
 - O_EXCL
 - When used with O_CREAT, if the file already exists it is an error and the open will fail

open()

- flags (cont'd)
 - **O_NOCTTY**
 - If pathname refers to a terminal device it will not become the process's controlling terminal even if the process does not have one
 - **O_TRUNC**
 - If the file already exists it will be truncated
 - **O_APPEND**
 - Initially, and before each write, the file pointer is positioned at the end of the file

open()

- flags (cont'd)
 - **O_NONBLOCK (O_NDELAY)**
 - Neither the open nor any subsequent operations on the file descriptor which is returned will cause the calling process to wait
 - regular file: mandatory lock
 - FIFO: open, read, write
 - **O_SYNC**
 - Any writes on the resulting file descriptor will block the calling process until the data has been physically written to the underlying hardware

open()

- flags (cont'd)
 - **O_NOFOLLOW**
 - If pathname is a symbolic link, then the open fails
 - define _GNU_SOURCE before including <fcntl.h>
 - **O_DIRECTORY**
 - If pathname is not a directory, cause the open to fail
 - define _GNU_SOURCE before including <fcntl.h>
 - Directory can be opened with open(), but can't read by read() system call. Should use readdir() system call.

open()

- flags (cont'd)
 - O_LARGE
 - On 32-bit systems that support the Large Files System, allow files whose sizes cannot be represented in 31 bits to be opened.
 - define _GNU_SOURCE before including <fcntl.h>
 - cf) open64, lseek64

open()

- mode
 - specifies the permissions to use if a new file is created
 - modified by the process's umask
 - the permissions of the created file are (mode & ~umask)
 - should always be specified when O_CREAT is in the flags, and is ignored otherwise
- return value
 - return the new file descriptor, or -1 if an error occurred

open()

■ example

```
int fd;  
fd = open("/etc/passwd", O_RDONLY);  
fd = open("/etc/passwd", O_RDWR);  
  
fd = open("ap", O_RDWR | O_APPEND);  
fd = open("ap", O_RDWR | O_CREAT | O_EXCL, 0644);
```

open()

■ example (synchronization)

```
#include <stdio.h>
#include <unistd.h>
#include <fcntl.h>
#include <errno.h>

#define LOCKFILE "lockfile"
#define DELAY 10000000

void delay(void)
{
    int i;
    for (i = 0; i < DELAY; i++);
}
```

open()

```
int main(void)
{
    int fd, i;

    while ((fd = open(LOCKFILE, O_WRONLY | O_CREAT | O_EXCL, 0644)) < 0) {
        if (errno != EEXIST) {
            perror("open");
            exit(1);
        }
    }
    for (i = 'a'; i <= 'z'; i++) {
        putchar(i);
        fflush(stdout);
        delay();
    }
    close(fd);
    unlink(LOCKFILE);

    return 0;
}
```

creat()

❑ synopsis

- `#include <sys/types.h>`
- `#include <sys/stat.h>`
- `#include <fcntl.h>`
- `int creat(const char *pathname, mode_t mode);`

❑ description

- create a new file
- equivalent to open with flags equal to
`O_CREAT|O_WRONLY|O_TRUNC`

close()

❑ synopsis

- `#include <unistd.h>`
- `int close(int fd);`

❑ description

- closes a file descriptor
- if the descriptor was the last reference to a file which has been removed using “unlink” the file is deleted

close()

- when a process terminates, all open files are automatically closed by the kernel
- return value
 - zero on success, or -1 if an error occurred

close()

■ example

```
#include <unistd.h>
#include <fcntl.h>
#include <stdio.h>

int main(void)
{
    int fd, n;
    char buf[10];
    fd = open("testfile", O_RDWR | O_CREAT | O_TRUNC,
0644);
    if (fd < 0) {
        perror("open");
        exit(1);
    }
```

close()

```
system("ls testfile");

unlink("testfile");
system("ls testfile");

write(fd, "hello\n", 6);
lseek(fd, 0, SEEK_SET);
n = read(fd, buf, 10);
write(STDOUT_FILENO, buf, n);

close(fd);
}
```

lseek()

❑ synopsis

- `#include <sys/types.h>`
- `#include <unistd.h>`
- `off_t lseek(int fildes, off_t offset, int whence);`

❑ description

- repositions the offset of the file descriptor fildes to the argument offset

lseek()

- whence
 - SEEK_SET
 - The offset is set to offset bytes.
 - SEEK_CUR
 - The offset is set to its current location plus offset bytes.
 - SEEK_END
 - The offset is set to the size of the file plus offset bytes.

lseek()

- hole

- allows the file offset to be set beyond the end of the existing end-of-file of the file
 - If data is later written at this point, subsequent reads of the data in the gap return bytes of zeros

- return value

- success : the resulting offset location as measured in bytes from the beginning of the file
 - error : -1

lseek()

■ example

- off_t curpos;
- curpos = lseek(fd, 0, SEEK_CUR);
- lseek(fd, 0, SEEK_SET);
- lseek(fd, 0, SEEK_END);
- lseek(fd, -10, SEEK_CUR);
- lseek(fd, 100, SEEK_END);

lseek()

■ example

```
#include <unistd.h>
#include <fcntl.h>

int main(void)
{
    if (lseek(STDIN_FILENO, 0, SEEK_CUR) == -1)
        printf("cannot seek\n");
    else
        printf("seek OK\n");
    exit(0);
}
```

lseek()

■ result

```
$ gcc lseek.c
$ cat < /etc/motd | a.out
cannot seek
$ a.out < /etc/motd
seek OK
```

lseek()

example

```
#include <unistd.h>
#include <fcntl.h>

int main(void)
{
    int fd;

    fd = creat("holefile", 0644);
    write(fd, "hello", 5);

    lseek(fd, 10, SEEK_CUR);
    write(fd, "world", 5);

    lseek(fd, 8192, SEEK_SET);
    write(fd, "bye", 3);

    close(fd);
    return 0;
}
```

lseek()

■ result

```
$ ls -l holefile  
-rw-r--r-- 1 kim  stud      8195 Jul 18 21:37 holefile  
$ od -c holefile  
0000000  h  e  l  l  o  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  w  
0000020  o  r  l  d  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  
0000040  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  \0  
*  
0020000  b  y  e  
0020003  
$ du holefile  
8  holefile
```

read()

❑ synopsis

- `#include <unistd.h>`
- `ssize_t read(int fd, void *buf, size_t count);`

❑ description

- attempts to read up to count bytes from file descriptor fd into the buffer starting at buf

read()

- If count is zero, read() returns zero and has no other results
- return value
 - On success, the number of bytes read
 - zero indicates end of file
 - On error, -1 is returned

write()

❑ synopsis

- `#include <unistd.h>`
- `ssize_t write(int fd, const void *buf, size_t count);`

❑ description

- writes up to count bytes to the file referenced by the file descriptor fd from the buffer starting at buf

write()

- return value
 - the number of bytes written
 - zero indicates nothing was written
 - On error, -1

read() & write()

```
#include <unistd.h>
#define BUFFSIZE 8192
int main(void)
{
    int n;
    char buf[BUFFSIZE];

    while ((n=read(STDIN_FILENO,buf,BUFFSIZE))>0)
        if (write(STDOUT_FILENO,buf,n)!=n)
            printf("write error\n");

    if (n<0)
        printf("read error\n");
    exit(0);
}
```

Atomic Operations

❑ appending to a file

```
lseek(fd, 0, SEEK_END);
write(fd, buf, 256);
```

❑ creating a file

```
if ((fd = open(pathname, O_WRONLY)) < 0) {
    if (errno == ENOENT)
        fd = creat(pathname, 0644);
    else {
        perror("open");
        exit(1);
    }
}
```

dup() and dup2()

❑ synopsis

- `#include <unistd.h>`
- `int dup(int oldfd);`
- `int dup2(int oldfd, int newfd);`

❑ description

- create a copy of the file descriptor `oldfd` and returns a new file descriptor
- `dup2` makes `newfd` be the copy of `oldfd`, closing `newfd` first if necessary.

dup() and dup2()

- close_on_exec flag is also copied.
- The old and new descriptors may be used interchangeably
 - if the file position is modified by using lseek on one of the descriptors the position is also changed for the other
- two descriptors do not share the close-on-exec flag
- dup uses the lowest-numbered unused descriptor for the new descriptor

dup() and dup2()


- return value
 - the new descriptor, or -1 if an error occurred

❑ example

- 1) int fd = dup(STDOUT_FILENO);
- 2) fd = open("input_file", O_RDONLY);
 dup2(fd, STDIN_FILENO);
- 3) fd = open("input_file", O_RDONLY);
 close(STDIN_FILENO);
 dup(fd);

dup() and dup2()

- fd = open("/etc/passwd", O_RDONLY);
- newfd = dup(fd);




dup() and dup2()

▣ 2개의 다른 process가 같은 파일을 open

file descriptor
table

file table

inode table



dup() and dup2()

Example

```
#include <unistd.h>
#include <fcntl.h>

int main(void)
{
    int fd;
    fd = creat("dup_result", 0644);
    dup2(fd, STDOUT_FILENO);
    close(fd);
    printf("hello world\n");
    return 0;
}
```

```
$ cat dup_result
hello world
```

fcntl()

■ synopsis

- `#include <unistd.h>`
- `#include <fcntl.h>`

- `int fcntl(int fd, int cmd);`
- `int fcntl(int fd, int cmd, long arg);`

fcntl()

■ Description

- 이미 열려진 파일에 대해서 현재의 상태를 읽어 오거나 새로운 상태를 설정한다.
- manipulate file descriptor
- performs one of various miscellaneous operations on fd
- The operation in question is determined by cmd

fcntl()

- cmd

- F_DUPFD
 - Makes arg be a copy of fd
- F_GETFD
 - Read the close-on-exec flag
 - If the low-order bit is 0, the file will remain open across exec
 - otherwise it will be closed.
- F_SETFD
 - Set the close-on-exec flag to the value specified by arg
 - only the least significant bit is used
- F_GETFL
 - Read the descriptor's flags
 - all flags (as set by open()) are returned

fcntl()

- example

```
int flag = fcntl(fd, F_GETFL);
if ((flag & O_ACCMODE) == O_RDONLY)
    printf("opened with RONLY");
else if ((flag & O_ACCMODE) == O_WRONLY)
    printf("opened with WRONLY");
else
    printf("opened with RDWR");
if (flag & O_APPEND)
    printf(",APPEND\n");
else
    printf("\n");
```

fcntl()

- F_SETFL
 - Set the descriptor's flags to the value specified by arg
 - Only O_APPEND, O_NONBLOCK and O_ASYNC may be set
 - the other flags are unaffected
- example
 - int flag = fcntl(fd, F_GETFL);
 - fcntl(fd, F_SETFL, flag | O_APPEND);

fcntl()

- return value
 - F_DUPFD
 - The new descriptor
 - F_GETFD
 - the close-on-exec flag
 - F_GETFL
 - value of descriptor's flags
 - F_SETFD, F_SETFL
 - On success, 0
 - On error, -1

/dev/fd

■ /dev/fd

- opening the file /dev/fd/n is equivalent to duplicating descriptor n
- example
- `fd = open("/dev/fd/0", O_RDONLY);`
- `fd = dup(0); /* STDIN_FILENO */`

Example - copy

```
#include <sys/types.h>
#include <fcntl.h>
#include <unistd.h>
#include <stdio.h>

main(int argc, char *argv[]);
{
    int fd1, fd2, n;
    char buf[1024];
    if (argc < 3) {
        fprintf(stderr, "Usage: %s src dest\n", argv[0]);
        exit(1);
}
```

Example - copy

```
if ((fd1 = open(argv[1], O_RDONLY)) < 0) {
    perror("Error:");
    exit(1);
}
if ((fd2 = creat(argv[2], 0644)) < 0) {
    perror("Error:");
    exit(1);
}
while ((n = read(fd1, buf, 1024)) > 0)
    write(fd2, buf, n);
close(fd1);
close(fd2);
}
```

Example – large file

```
#include <unistd.h>
#include <fcntl.h>
#include <sys/types.h>
main()
{
    int fd;
    loff_t p = 1LL << 32;
    if ((fd = open64("large", O_CREAT|O_RDWR|O_TRUNC, 0644)) < 0)
    {
        perror("open");
        exit(1);
    }
```

Example – large file

```
if (lseek64(fd, p, SEEK_SET) < 0) {  
    perror("llseek");  
    exit(0);  
}  
write(fd, "hello", 5);  
close(fd);  
}
```