

# Files and Directories

# stat( ), fstat( ), lstat( )

## ■ synopsis

- `#include <sys/stat.h>`
- `#include <unistd.h>`
- `int stat(const char *file_name, struct stat *buf);`
- `int fstat(int filedes, struct stat *buf);`
- `int lstat(const char *file_name, struct stat *buf);`

# stat( ), fstat( ), lstat( )

## ▣ stat

- pathname이 가리키는 파일에 관한 정보를 buf가 가리키는 struct stat에 저장한다.

## ▣ fstat

- stat과 동일하지만 파일을 이미 open한 파일 descriptor로 지정한다.

## ▣ lstat

- stat과 유사하지만 파일이 symbolic link일 경우 symbolic link가 참조하는 파일이 아니라 symbolic link 그 자체의 파일에 대한 정보를 return한다

# stat( ), fstat( ), lstat( )

```
struct stat {  
    dev_t st_dev; /* device */  
    ino_t st_ino; /* inode */  
    mode_t st_mode; /* protection */  
    nlink_t st_nlink; /* number of hard links */  
    uid_t st_uid; /* user ID of owner */  
    gid_t st_gid; /* group ID of owner */  
    dev_t st_rdev; /* device type (if inode device) */  
    off_t st_size; /* total size, in bytes */  
    unsigned long st_blksize; /* blocksize for I/O */  
    unsigned long st_blocks; /* # of blocks allocated */  
    time_t st_atime; /* time of last access */  
    time_t st_mtime; /* time of last modification */  
    time_t st_ctime; /* time of last change */  
};
```

# stat( ), fstat( ), lstat( )

## ■ description

- return information about the specified file
- do not need any access rights to the file
- need search rights to all directories named in the path leading to the file

## ■ all return a stat structure

- st\_blocks
  - the size of the file in 512-byte blocks
- st\_blksize
  - the "preferred" block size for efficient file system I/O
  - Writing to a file in smaller chunks may cause an inefficient read-modify-rewrite

# stat( ), fstat( ), lstat( )

- st\_atime
  - changed by exec(), mknod(), pipe(), utime(), read()
  - atime may not be updated
    - ex) mounting with noatime option
- st\_mtime
  - changed by mknod(), truncated(), utime(), write()
  - st\_mtime of a directory is changed by creation or deletion of files in that directory.
- st\_ctime
  - changed by writing or by setting inode information
  - owner, group, link count, mode, etc.

# stat( ), fstat( ), lstat( )

## ▣ st\_mode:

- st\_mode는 아래 그림과 같이 file type, special bit, file permission bit의 3 부분으로 나눈다

type	special	permission
------	---------	------------

15

12

9

0

# stat( ), fstat( ), lstat( )

## file type

- #define S\_IFMT 00170000 /\* type of file \*/
- #define S\_IFSOCK 0140000 /\* socket \*/
- #define S\_IFLNK 0120000 /\* symbolic link \*/
- #define S\_IFREG 0100000 /\* regular \*/
- #define S\_IFBLK 0060000 /\* block special \*/
- #define S\_IFDIR 0040000 /\* directory \*/
- #define S\_IFCHR 0020000 /\* character special \*/
- #define S\_IFIFO 0010000 /\* FIFO \*/

# stat( ), fstat( ), lstat( )

## ■ file type macro

Macro	Type of file
S_ISREG()	regular file
S_ISDIR()	directory file
S_ISCHR()	character special file
S_ISBLK()	block special file
S_ISFIFO()	pipe or FIFO
S_ISLNK()	symbolic link
S_ISSOCK()	socket

## ■ special bits

- #define S\_ISUID 0004000 /\* set uid on execution \*/
- #define S\_ISGID 0002000 /\* set group id on execution \*/
- #define S\_ISVTX 0001000 /\* save text(sticky bit) \*/

# stat( ), fstat( ), lstat( )

## ■ permission bits

- #define S\_IRUSR 00400
- #define S\_IWUSR 00200
- #define S\_IXUSR 00100
- #define S\_IRGRP 00040
- #define S\_IWGRP 00020
- #define S\_IXGRP 00010
- #define S\_IROTH 00004
- #define S\_IWOTH 00002
- #define S\_IXOTH 00001

```
/* read permission: owner */  
/* write permission: owner */  
/* execute permission: owner */  
/* read permission: group */  
/* write permission: group */  
/* execute permission: group */  
/* read permission: other */  
/* write permission: other */  
/* execute permission: other */
```

# stat( ), fstat( ), lstat( )

- return value

- On success, zero is returned
- On error, -1 is returned

- cf) “stat” shell command

- example

```
struct stat buf;  
stat("/etc/passwd", &buf);  
if (buf.st_mode & S_IRUSR) {
```

```
...
```

# statfs(), fstatfs()

## ■ synopsis

- `#include <sys/vfs.h>`
- `int statfs(const char *path, struct statfs *buf);`
- `int fstatfs(int fs, struct statfs *buf);`

## ■ description

- returns information about a mounted file system
- path is the path name of any file within the mounted file system

# statfs(), fstatfs()

```
struct statfs {  
    long  f_type; /* type of filesystem (magic number, EXT2: 0xEF53) */  
    long  f_bsize; /* optimal transfer block size */  
    long  f_blocks; /* total data blocks in file system */  
    long  f_bfree; /* free blocks in fs */  
    long  f_bavail; /* free blocks avail to non-superuser */  
    long  f_files; /* total file nodes in file system */  
    long  f_ffree; /* free file nodes in fs */  
    fsid_t f_fsid; /* file system id */  
    long  f_namelen; /* maximum length of filenames */  
    long  f_spare[6]; /* spare for later */  
};
```

# File Access Permissions

## ■ real ID and effective ID

- real user(group) ID
  - the user(group) ID who executed the process
- effective user(group) ID
  - user(group) ID used to assign ownership of newly created file or to check access permission when system call is done
- real user(group) ID can be different from the real user(group), in the case of programs with the setuid(setgid) bit set

# File Access Permissions

## example

```
-rwxr-xr-x student bitgrp 43753 a.out
```

```
chmod u+s a.out
```

```
-rwsr-xr-x student bitgrp 43753 a.out
```

- the effective user id of the process which is executing “a.out” is “student”, and the real user id is “guest”

# File Access Permissions

## ■ regular file

- read permission
  - can open with O\_RDONLY and O\_RDWR flags
- write permission
  - can open with O\_WRONLY and O\_RDWR flags
- must have write permission to specify the O\_TRUNC

## ■ example

```
user = stud, group = bitgrp
--wxr-x--- stud bitgrp testfile
```

# File Access Permissions

## ■ directory

- whenever want to open any type of file by name, must have execute permission in each directory mentioned in the name
- read permission for directory
  - lets us read the directory, obtaining a list of all the filenames in the directory
- execution permission for directory
  - lets us pass through the directory when it is a component of a pathname that we are trying to access

# File Access Permissions

---

- cannot create a new file in a directory unless we have write permission and execute permission in the directory
  - to delete a file, we need write permission and execute permission in the directory containing the file
    - do not need read permission or write permission for the file itself
- ❑ ownership of new files
- the effective user ID and group ID of the creating process

# access( )

## ■ synopsis

- `#include <unistd.h>`
- `int access(const char *pathname, int mode);`

## ■ description

- checks whether the process would be allowed to read, write or test for existence of the file whose name is pathname
- If pathname is a symbolic link permissions of the file referred to by this symbolic link are tested

# access( )

- mode
  - R\_OK
    - file exist, read permission
  - W\_OK
    - file exist, write permission
  - X\_OK
    - file exist, execution permission
  - F\_OK
    - file exist
- The tests depend on the permissions of the directories occurring in the path to the file

# access( )

- The check is done with the process's real uid and gid
  - This is to allow set-UID programs to easily determine the invoking user's authority
- return value
  - On success zero is returned
  - On error -1 is returned
  - returns an error if any of the access types in the requested call fails, even if other types might be successful
  - `access("bit", R_OK | W_OK)` : check read and write permission

# umask( )

## ■ synopsis

- #include <sys/types.h>
- #include <sys/stat.h>
- mode\_t umask(mode\_t mask);

## ■ description

- set file creation mask
  - sets the umask to mask & 0777
- permissions in the umask are turned off from the mode argument to open()

# umask( )

- return value
  - This system call always succeeds and the previous value of the mask is returned
- cf) “umask” shell command
- example

```
umask(022);
fd = creat("tmp", 0666);
...
● permission of "tmp" is -rw-r--r-- (0644)
```

# chmod( ) and fchmod( )

## ■ synopsis

- `#include <sys/types.h>`
- `#include <sys/stat.h>`
- `int chmod(const char *path, mode_t mode);`
- `int fchmod(int fildes, mode_t mode);`

## ■ description

- change permissions of a file
- The mode of the file given by path or referenced by fildes is changed

# chmod( ) and fchmod( )

- mode is specified by or'ing the following
  - S\_ISUID, S\_ISGID, S\_ISVTX,  
S\_I{R,W,X}{USR,GRP,OTH}, S\_IRWX{U,G,O}
- effective UID of the process must be zero or must match the owner of the file
- SGID bit is turned off when
  - effective UID of the process is not zero
  - the group of the file does not match the effective group ID of the process or one of its supplementary group IDs

# chmod( ) and fchmod( )

- return value
  - On success, zero is returned
  - On error, -1 is returned

# Sticky bit

## ■ S\_ISVTX bit

- saved text
- set for an executable program file
  - the first time the program was executed a copy of the program's text was saved in the swap area when the process terminated
  - causes the program to load into memory faster
  - today's Unix systems have a virtual memory system and a faster file system, the need for this technique has disappeared

# Sticky bit

- set for a directory
  - a file in the directory can be removed or renamed only if
    - the user has write permission for the directory and
    - the user owns the file or owns the directory or the user is superuser

## example

- `chmod("/tmp", S_ISVTX | 0777);`
- permission of “/tmp” is drwxrwxrwt

# chown( ), fchown( ), lchown( )

## ■ synopsis

- #include <sys/types.h>
- #include <unistd.h>
  
- int chown(const char \*path, uid\_t owner, gid\_t group);
- int fchown(int fd, uid\_t owner, gid\_t group);
- int lchown(const char \*path, uid\_t owner, gid\_t group);

# chown( ), fchown( ), lchown( )

## ■ description

- The owner of the file specified by path or by fd is changed
- Only the super-user may change the owner of a file
- The owner of a file may change the group of the file to any group of which that owner is a member
- when the group of an executable file is changed by a non-super-user, setgid bit is cleared.

# chown( ), fchown( ), lchown( )

- If the owner or group is specified as -1, then that ID is not changed
- return value
  - On success, zero is returned
  - On error, -1 is returned

# truncate( ) and ftruncate( )

## ■ synopsis

- #include <unistd.h>
- int truncate(const char \*path, off\_t length);
- int ftruncate(int fd, off\_t length);

## ■ description

- truncate a file to a specified length
- If the file previously was larger than length, the extra data is lost
- If the file previously was smaller than length, the hole is created

# truncate( ) and ftruncate( )

- With ftruncate, the file must be open for writing
- return value
  - On success, zero is returned
  - On error, -1 is returned

# link( )

## ■ synopsis

- `#include <unistd.h>`
- `int link(const char *oldpath, const char *newpath);`

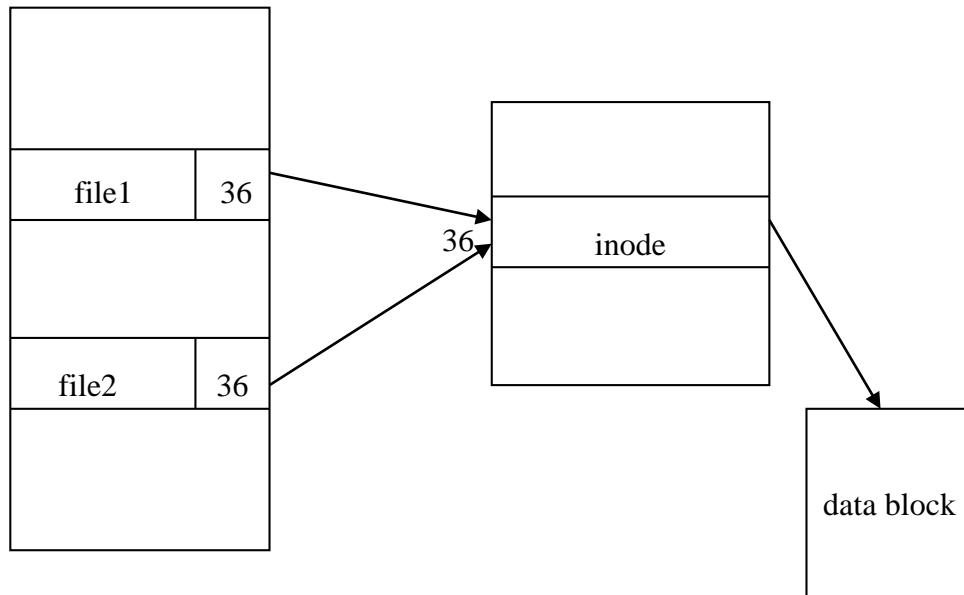
## ■ description

- make a new name for a file
  - creates a new link (also known as a hard link) to an existing file
- If newpath exists it will not be overwritten

# link( )

- This new name may be used exactly as the old one for any operation
  - both names refer to the same file
  - have the same permissions and ownership
  - it is impossible to tell which name was the original
- Hard links, as created by link, cannot span file systems
- Cannot create a link to a directory

# link( )



link("file1", "file2");

# link( )

## ■ link count

- regular file
  - the link count is the number of the names
- directory
  - all the directories have at least two links (“.” and directory name)
  - all the directories have a link to the parent directory
  - the link count of a directory is  $2 + \#$  of sub directories

# unlink( )

## ■ synopsis

- `#include <unistd.h>`
- `int unlink(const char *pathname);`

## ■ description

- delete a name and possibly the file it refers to
  - deletes a name from the file system
  - If that name was the last link to a file and no processes have the file open the file is deleted

# unlink( )

- If the name was the last link to a file but any processes still have the file open the file will remain in existence until the last file descriptor referring to it is closed
- If the name referred to a symbolic link the link is removed
- return value
  - On success, zero is returned. On error, -1 is returned

# remove( )

## ■ synopsis

- `#include <stdio.h>`
- `int remove(const char *pathname);`

## ■ description

- library function(not system call)
- delete a name and possibly the file it refers to
  - It calls unlink for files, and rmdir for directories
- return value
  - On success, zero is returned. On error, -1 is returned

# rename( )

## ■ synopsis

- `#include <stdio.h>`
- `int rename(const char *oldpath, const char *newpath);`

## ■ description

- renames a file or a directory, moving it between directories if required
- oldpath and newpath should be in the same file system

# rename( )

- If newpath already exists it will be atomically replaced
  - if oldpath is not a directory, newpath should not be a directory
  - if oldpath is a directory and newpath exists, newpath should be an empty directory
- If oldpath refers to a symbolic link the link is renamed
- if newpath refers to a symbolic link the link will be overwritten

# Symbolic Link

- symbolic link is indirect pointer to a file
  - hard link points directly to the inode of the file
  - limitation of hard links
    - require that the link and the file reside in the same file system
    - a hard link to a directory can't be created
  - there are no file system limitations on a symbolic link
  - anyone can create a symbolic link to a directory

# symlink( )

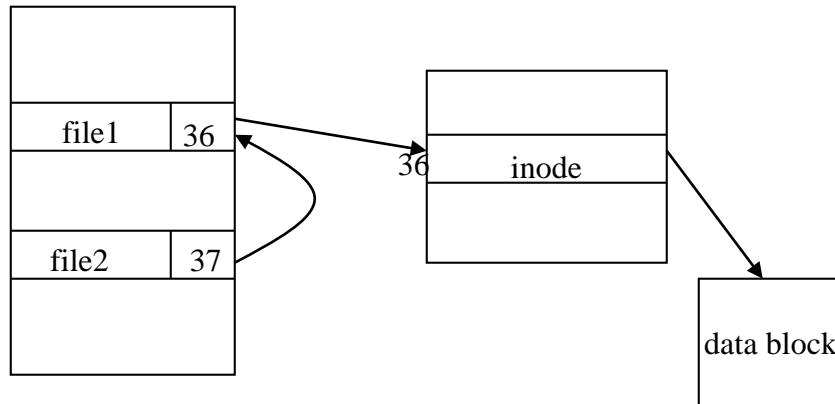
## ■ synopsis

- `#include <unistd.h>`
- `int symlink(const char *oldpath, const char *newpath);`

## ■ description

- creates a symbolic link named newpath which contains the string oldpath
- Symbolic links are interpreted at run-time
- dangling link
  - may point to an non-existing file

# symlink( )



```
$ ln -s testfile newfile
```

```
$ ls -l newfile
```

```
1 lrwxrwxrwx 1 yhshin users 8 Aug 27 20:02
```

```
newfile -> testfile
```

```
$ rm testfile
```

```
$ cat newfile
```

```
cat: newfile: No such file or directory
```

하드링크와 달리 원래 파일을 지우면 reference 할 수 없다.

# `symlink( )`

- The permissions of a symbolic link are irrelevant
  - the ownership is ignored when following the link
  - permission is checked when removal or renaming of the link is requested and the link is in a directory with the sticky bit set
- If newpath exists it will not be overwritten
- return value
  - On success, zero is returned. On error, -1 is returned

# readlink( )

## ■ synopsis

- `#include <unistd.h>`
- `int readlink(const char *path, char *buf, size_t bufsiz);`

## ■ description

- read value of a symbolic link
  - places the contents of the symbolic link path in the buffer buf, which has size bufsiz
  - does not append a NULL character to buf

# readlink( )

- return value

- the count of characters placed in the buffer if it succeeds
  - -1 if an error occurs

# utime( )

## ❑ synopsis

- `#include <sys/types.h>`
- `#include <utime.h>`
- `int utime(const char *filename, struct utimbuf *buf);`

## ❑ description

- change access and/or modification times of an inode
  - changes the access and modification times of the inode specified by filename to the actime and modtime fields of buf respectively

# utime( )

- If buf is NULL, then the access and modification times of the file are set to the current time

```
struct utimbuf {  
    time_t actime; /* access time */  
    time_t modtime; /* modification time */  
};
```

- return value
  - On success, zero is returned. On error, -1 is returned

# utime( ) - example

```
#include<sys/types.h>
#include<sys/stat.h>
#include<fcntl.h>
#include<utime.h>
#include <stdio.h>
#include<unistd.h>

int main(int argc, char *argv[])
{
    int    i;
    struct stat    statbuf;
    struct utimbuf        timebuf;
```

# utime( ) - example

```
for (i = 1; i < argc; i++) {
    if (stat(argv[i], &statbuf) < 0) {
        fprintf(stderr, "%s: stat error", argv[i]);
        continue;
    }
    if (open(argv[i], O_RDWR | O_TRUNC) < 0) {
        fprintf(stderr, "%s: open error", argv[i]);
        continue;
    }
}
```

# utime( ) - example

```
timebuf.actime = statbuf.st_atime;
timebuf.modtime = statbuf.st_mtime;
if (utime(argv[i], &timebuf) < 0) {
    fprintf(stderr, "%s: utime error", argv[i]);
    continue;
}
}
```

# mkdir( )

## ■ synopsis

- `#include <sys/stat.h>`
- `#include <sys/types.h>`
- `#include <fcntl.h>`
- `#include <unistd.h>`
- `int mkdir(const char *pathname, mode_t mode);`

# `mkdir( )`

## ■ description

- attempts to create a directory named pathname
- mode specifies the permissions to use
  - modified by the process's umask in the usual way
- ownership
  - The newly created directory will be owned by the effective uid of the process
  - If the directory containing the file has the set group id bit set, the new directory will inherit the group ownership from its parent

# **mkdir( )**

- otherwise it will be owned by the effective gid of the process
- set group id bit
  - If the parent directory has the set group id bit set then so will the newly created directory
- return value
  - zero on success, or -1 if an error occurred

# rmdir( )

## ■ synopsis

- `#include <unistd.h>`
- `int rmdir(const char *pathname);`

## ■ description

- deletes a directory, which must be empty
- return value
  - On success, zero is returned. On error, -1 is returned

# Reading Directories

## ■ synopsis

- #include <sys/types.h>
- #include <dirent.h>
- DIR \*opendir(const char \*name);
- struct dirent \*readdir(DIR \*dir);
- off\_t telldir(DIR \*dir);
- void seekdir(DIR \*dir, off\_t offset);
- void rewinddir(DIR \*dir);
- int closedir(DIR \*dir);

# Reading Directories

- `DIR *opendir(const char *name);`
  - opens a directory stream corresponding to the directory name, and returns a pointer to the directory stream
  - returns a pointer to the directory stream or NULL if an error occurred
- `struct dirent *readdir(DIR *dir);`
  - returns a pointer to a dirent structure representing the next directory entry in the directory stream pointed to be dir

# Reading Directories

- The data returned by readdir() is overwritten by subsequent calls to readdir() for the same directory stream
- returns a pointer to a dirent structure, or NULL if an error occurs or end-of-file is reached

```
struct dirent {  
    ino_t d_ino;  
    char d_name[NAMEMAX+1];  
}
```

# Reading Directories

- `off_t telldir(DIR *dir);`
  - returns the current location associated with the directory stream `dir`.
  - returns `-1` if an error occurs
- `void seekdir(DIR *dir, off_t offset);`
  - sets the location in the directory stream from which the next `readdir()` call will start.
  - should be used with an offset returned by `telldir()`.

# Reading Directories

- `void rewinddir(DIR *dir);`
  - resets the position of the directory stream dir to the beginning of the directory
- `int closedir(DIR *dir);`
  - closes the directory stream associated with dir
  - returns 0 on success or -1 on failure

# scandir( )

## ■ synopsis

- #include <dirent.h>
- int scandir(const char \*dir,  
              struct dirent \*\*\*namelist,  
              int (\*select)(const struct dirent \*),  
              int (\*compar)(const struct dirent \*\*,    const  
                  struct dirent \*\*));

## ■ return value

- the number of directory entries selected or -1  
(error)

# scandir( )

## ■ description

- scans the directory dir, calling select() on each directory entry.
- Entries for which select() returns non-zero are stored in strings allocated via malloc(), sorted using comparison function compar(), and collected in array namelist which is allocated via malloc().
- If select is NULL, all entries are selected.

# scandir( )

## ■ example

```
#include <dirent.h>
#include <string.h>
int selfile(const struct dirent *p)
{
    if (strstr(p->d_name, ".c"))
        return 1;
    return 0;
}
int mysort(const struct dirent **a, const struct dirent **b)
{
    return strlen((*a)->d_name) - strlen((*b)->d_name);
}
```

# scandir( )

```
main(int argc, char *argv[])
{
    struct dirent **namelist;
    int n, i;
    if ((n = scandir(argv[1], &namelist, selfile, mysort)) < 0) {
        perror("scandir");
        exit(1);
    }
    for (i = 0; i < n; i++) {
        printf("%s\n", namelist[i]->d_name);
    }
}
```

# Example

```
#include <sys/types.h>
#include <dirent.h>
#include <sys/stat.h>
#include <unistd.h>
#include <stdio.h>

void access_perm(char *perm, mode_t mode)
{
    int i;
    char permchar[] = "rwx";
    memset(perm, '-', 10);
    perm[10] = '\0';
```

# Example

```
if (S_ISDIR(mode)) perm[0] = 'd';
else if (S_ISCHR(mode)) perm[0] = 'c';
else if (S_ISBLK(mode)) perm[0] = 'b';
else if (S_ISFIFO(mode)) perm[0] = 'p';
else if (S_ISLNK(mode)) perm[0] = 'l';
for (i = 0; i < 9; i++) {
    if ((mode << i) & 0x100)
        perm[i+1] = permchar[i%3];
}
if (mode & S_ISUID) perm[3] = 's';
if (mode & S_ISGID) perm[6] = 's';
if (mode & S_ISVTX) perm[9] = 't';
}
```

# Example

```
main(int argc, char *argv[])
{
    DIR *dp;
    struct stat statbuf;
    struct dirent *dent;
    char perm[11];
    char pathname[80];

    if (argc < 2) exit(1);
    if (access(argv[1], F_OK) == -1) {
        perror("Error");
        exit(1);
    }
```

# Example

```
stat(argv[1], &statbuf);
if (!S_ISDIR(statbuf.st_mode)) {
    fprintf(stderr, "%s is not directory\n",
            argv[1]);
    exit(1);
}
if ((dp = opendir(argv[1])) == NULL) {
    perror("Error:");
    exit(1);
}
printf("Lists of Directory(%s):\n", argv[1]);
```

# Example

```
while((dent = readdir(dp)) != NULL) {  
    sprintf(pathname, "%s/%s", argv[1],  
            dent->d_name);  
    stat(pathname, &statbuf);  
    access_perm(perm, statbuf.st_mode);  
    printf("%s %8ld %s\n", perm, statbuf.st_size,  
          dent->d_name);  
}  
closedir(dp);  
}
```

# chdir( ) and fchdir( )

## ■ synopsis

- `#include <unistd.h>`
- `int chdir(const char *path);`
- `int fchdir(int fd);`

## ■ description

- changes the current directory to that specified in path
- fchdir uses an open file descriptor as an argument

# chdir( ) and fchdir( )

## example

```
#include <unistd.h>
#include <stdio.h>
main(void)
{
    if (chdir("/tmp") < 0) {
        perror("chdir");
        exit(1);
    }
    printf("chdir to /tmp succeeded\n");
}
```

# getcwd( )

## ■ synopsis

- `#include <unistd.h>`
- `char *getcwd(char *buf, size_t size);`

## ■ description

- Get current working directory
  - copies the absolute pathname of the current working directory to the array pointed to by buf, which is of length size

# getcwd( )

- If the current absolute path name would require a buffer longer than size elements, NULL is returned
- allocates the buffer dynamically using malloc() if buf is NULL on call
- return value
  - NULL on failure
  - buf on success

# chroot( )

## ■ synopsis

- `#include <unistd.h>`
- `int chroot(const char *path);`

## ■ description

- changes the root directory to that specified in path.
- only the super-user may change the root directory.

## ■ return value

- On success, zero is returned.
- On error, -1 is returned.

# sync( )

## ❑ synopsis

- `#include <unistd.h>`
- `int sync(void);`

## ❑ description

- first commits inodes to buffers, and then buffers to disk.

## ❑ return value

- always returns 0.

# fsync( )

## ■ synopsis

- `#include <unistd.h>`
- `int fsync(int fd);`

## ■ description

- copies all in-core parts of a file to disk.

## ■ return value

- On success, zero is returned.
- On error, -1 is returned

# fdatasync( )

## ■ synopsis

- `#include <unistd.h>`
- `int fdatasync(int fd);`

## ■ description

- flushes all data buffers of a file to disk.
- does not update the metadata.

## ■ return value

- On success, zero is returned.
- On error, -1 is returned