

Standard I/O Library

Streams and FILE objects

■ stream

- stream of bytes
- logical structure used to access physical file
- when we open or create a file with the standard I/O library, we say that we have associated a stream with the file

■ FILE object

- when we open a stream, fopen returns a pointer to a FILE object

Streams and FILE objects

- contains all the information required by the standard I/O library to manage the stream
 - file descriptor
 - pointer to a buffer for the stream
 - size of the buffer
 - counter of the number of characters in the buffer
 - error flag
- predefined stream
 - stdin, stdout, stderr (<stdio.h>)

Buffering

■ goal

- use the minimum number of read and write system calls

■ three types of buffering

● fully buffered

- actual I/O takes place when the standard I/O buffer is filled
- files that reside on disk are normally fully buffered

Buffering

- line buffered
 - performs I/O when a newline character('n') is encountered on input or output
 - used on a stream when it refers to a terminal (standard input and standard output)
 - actual I/O might take place if we fill the buffer before writing a newline
 - whenever input is requested through the standard I/O library from unbuffered stream or from a line-buffered stream, all line-buffered output streams are flushed

Buffering

- unbuffered
 - does not buffer the characters
 - ex) standard error stream

stream buffering operations

■ synopsis

- #include <stdio.h>
- int setbuf(FILE *stream, char *buf);
- int setbuffer(FILE *stream, char *buf, size_t size);
- int setlinebuf(FILE *stream);
- int setvbuf(FILE *stream, char *buf, int mode , size_t size);
- int fflush(FILE *stream);

stream buffering operations

■ description

- `int setvbuf(FILE *stream, char *buf, int mode , size_t size);`
 - may be used at any time on any open stream to change its buffer
 - mode parameter
 - `_IONBF` : unbuffered
 - `_IOLBF` : line buffered
 - `_IOFBF` : fully buffered

stream buffering operations

- buf parameter
 - should point to a buffer at least size bytes long
(except for unbuffered files)
 - buf will be used instead of the current buffer
 - if NULL, only the mode is affected
- a new buffer will be allocated on the next read or write operation
- can only change the mode of a stream when it is not “active”
 - before any I/O
 - immediately after a call to fflush()

stream buffering operations

- `int setbuf(FILE *stream, char *buf);`
 - `setvbuf(stream, buf, buf ? _IOFBF : _IONBF, BUFSIZ);`
 - `BUFSIZ`
 - default buffer size(8192 bytes)
- `int setbuffer(FILE *stream, char *buf, size_t size);`
 - the same as `setbuf()`, except that the size of the buffer is up to the caller
- `int setlinebuf(FILE *stream);`
 - `setvbuf(stream, (char *)NULL, _IOLBF, 0);`

stream buffering operations

- `fflush()`

- forces a write of all buffered data for the given output or update stream via the stream's underlying write function
- If the stream argument is NULL, `fflush` flushes all open output streams
- return value
 - Upon successful completion 0 is returned
 - Otherwise, EOF is returned

opening a stream

■ synopsis

- `#include <stdio.h>`
- `FILE *fopen (const char *path, const char *mode);`
- `FILE *fdopen (int fildes, const char *mode);`
- `FILE *freopen (const char *path, const char *mode, FILE *stream);`

opening a stream

■ description

- FILE *fopen (const char *path, const char *mode);
 - opens the file whose name is the string pointed to by path and associates a stream with it
 - mode
 - r : open file for reading
 - r+ : open for reading and writing
 - w : truncate file to zero length or create file for writing

opening a stream

- w+ : open for reading and writing. the file is created if it does not exist, otherwise it is truncated
- a : open for writing. The file is created if it does not exist. The stream is positioned at the end of the file
- a+ : open for reading and writing. The file is created if it does not exist. The stream is positioned at the end of the file
- Any created files will have mode 0666, as modified by the process' umask value

opening a stream

- FILE *fdopen (int fildes, const char *mode);
 - associates a stream with the existing file descriptor, fildes
- FILE *fopen (const char *path, const char *mode, FILE *stream);
 - opens the file whose name is the string pointed to by path and associates the stream pointed to by stream with it.
 - The original stream (if it exists) is closed.
 - The primary use of the fopen function is to change the file associated with a standard text stream (stderr, stdin, or stdout)

opening a stream

- return value

- Upon successful completion fopen, fdopen and freopen return a FILE pointer. Otherwise, NULL is returned

Restriction	r	w	a	r+	w+	a+
file must already exist	•			•		
previous contents of file discarded		•			•	
stream can be read	•			•	•	•
stream can be written		•		•	•	•
stream can be written only at end			•			•

closing a stream

❑ synopsis

- `#include <stdio.h>`
- `int fclose(FILE *stream);`

❑ description

- dissociates the named stream from its underlying file or set of functions
- if the stream was being used for output, any buffered data is written first, using `fflush()`

reading and writing a stream

■ type of I/O operation

- character-at-a-time I/O (fgetc, fputc)
- line-at-a-time I/O (fgets, fputs)
- direct I/O (binary I/O, fread, fwrite)
- formatted I/O (fprintf, fscanf)

character-at-a-time I/O

■ synopsis

- #include <stdio.h>
- int fgetc(FILE *stream);
- int getc(FILE *stream);
- int getchar(void);
- int fputc(int c, FILE *stream);
- int putc(int c, FILE *stream);
- int putchar(int c);
- int ungetc(int c, FILE *stream);

character-at-a-time I/O

■ description

- `int fgetc(FILE *stream);`
 - reads the next character from stream and returns it as an unsigned char cast to an int
 - returns EOF on end of file or error
- `int getc(FILE *stream);`
 - equivalent to fgetc() except that it may be implemented as a macro
- `int getchar(void);`
 - equivalent to getc(stdin)

character-at-a-time I/O

- `int fputc(int c, FILE *stream);`
 - writes the character c, cast to an unsigned char, to stream
- `int putc(int c, FILE *stream);`
 - equivalent to `fputc()` except that it may be implemented as a macro
- `int putchar(int c);`
 - equivalent to `putc(c,stdout)`
- return value
 - `fputc()`, `putc()`, `putchar()` return the character written as an unsigned char cast to an int or EOF on error

character-at-a-time I/O

- `int ungetc(int c, FILE *stream);`
 - pushes c back to stream, cast to unsigned char
 - it is available for subsequent read operation
 - only a single character can be pushed back
 - the character that we push back does not have to be the same character that was read
 - returns c on success, or EOF on error
- cf) `int getw(FILE *stream);`
`int putw(int c, FILE *stream);`
 - the value returned on error is also a legitimate data value.

ferror(), feof(), clearerr()

■ synopsis

- #include <stdio.h>
- void clearerr(FILE *stream);
- int feof(FILE *stream);
- int ferror(FILE *stream);

ferror(), feof(), clearerr()

■ description

- clearerr()

- clears the end-of-file and error indicators for the stream pointed to by stream

- feof()

- tests the end-of-file indicator for the stream
 - return nonzero if it is set

- ferror()

- tests the error indicator for the stream
 - returning non-zero if it is set

line-at-a-time I/O

■ synopsis

- `#include <stdio.h>`
- `char *fgets(char *s, int size, FILE *stream);`
- `char *gets(char *s);`
- `int fputs(const char *s, FILE *stream);`
- `int puts(const char *s);`

line-at-a-time I/O

■ description

- `gets()`

- reads a line from `stdin` into the buffer pointed to by `s` until either a terminating newline or EOF, which it replaces with '`\0`'.
 - No check for buffer overrun is performed

- `fgets()`

- reads in at most one less than `size` characters from stream and stores them into the buffer pointed to by `s`

line-at-a-time I/O

- Reading stops after an EOF or a newline. If a newline is read, it is stored into the buffer
 - '\0' is stored after the last character in the buffer
- `fputs()`
 - writes the string s to stream, without its trailing '\0'.
 - `puts()`
 - writes the string s and a trailing newline to stdout.

line-at-a-time I/O

- return value

- gets(), fgets()

- return s on success, and NULL on error or when end of file occurs while no characters have been read.

- puts(), fputs()

- return a non - negative number on success, or EOF on error.

binary I/O

■ synopsis

- `#include <stdio.h>`
- `size_t fread(void *ptr, size_t size, size_t nmemb, FILE *stream);`
- `size_t fwrite(const void *ptr, size_t size, size_t nmemb, FILE *stream);`

binary I/O

■ description

- `fread()`

- reads `nmemb` elements of data, each `size` bytes long, from the stream pointed to by `stream`, storing them at the location given by `ptr`

- `fwrite()`

- writes `nmemb` elements of data, each `size` bytes long, to the stream pointed to by `stream`, obtaining them from the location given by `ptr`

binary I/O

- return value

- return the number of items successfully read or written, not the number of characters
- If an error occurs, or the end-of-file is reached, the return value is a short item count (or zero).
- fread does not distinguish between end-of-file and error, and callers must use feof() and ferror() to determine which occurred.

binary I/O

example

```
struct rec tmp;  
struct rec item[10];
```

```
fread(&tmp, sizeof(struct rec), 1, fp);  
fwrite(item, sizeof(struct rec), 10, fp);
```

formatted I/O - output

❑ synopsis

- `#include <stdio.h>`
- `int printf(const char *format, ...);`
- `int fprintf(FILE *stream, const char *format, ...);`
- `int sprintf(char *str, const char *format, ...);`
- `int snprintf(char *str, size_t size, const char *format, ...);`

formatted I/O - output

■ description

- `printf()`

- write output to `stdout`, the standard output stream

- `fprintf()`

- write output to the given output stream

- `sprintf()`, `snprintf()`

- write to the character string `str`

- `snprintf()` does not write more than `size` bytes (including the trailing '`\0`')

formatted I/O - output

- return value

- These functions return the number of characters printed (not including the trailing '\0' used to end output to strings)
- negative value if output error
- snprintf() returns -1 if the output was truncated due to size limit

formatted I/O - input

■ synopsis

- `#include <stdio.h>`
- `int scanf(const char *format, ...);`
- `int fscanf(FILE *stream, const char *format, ...);`
- `int sscanf(const char *str, const char *format, ...);`

■ description

- `scanf()`
 - reads input from the standard input stream stdin

formatted I/O - input

- `fscanf()`
 - reads input from the stream pointer stream
- `sscanf()`
 - reads its input from the character string pointed to by str
- return value
 - return the number of input items assigned
 - EOF if input error or end of file before any conversion

positioning a stream

❑ synopsis

- `#include <stdio.h>`
- `int fseek(FILE *stream, long offset, int whence);`
- `long ftell(FILE *stream);`
- `void rewind(FILE *stream);`
- `int fgetpos(FILE *stream, fpos_t *pos);`
- `int fsetpos(FILE *stream, fpos_t *pos);`

positioning a stream

■ description

- `fseek()`

- sets the file position indicator for the stream pointed to by `stream`
- `SEEK_SET`, `SEEK_CUR`, `SEEK_END`
- successful call to the `fseek()` clears the end-of-file indicator for the stream and undoes any effects of the `ungetc()`
- return value
 - 0 if OK, -1 on error

positioning a stream

- `fseek()`
 - obtains the current value of the file position indicator for the stream
 - return value
 - current file position indicator if OK
 - -1L on error
- `rewind()`
 - `fseek(stream, 0, SEEK_SET); clearerr(stream);`

positioning a stream

- `fgetpos()`
 - store the current value of the file offset into the object referenced by pos
- `fsetpos()`
 - set the current value of the file offset into the object referenced by pos
- `return value`
 - 0 if OK, -1 on error

positioning a stream

example

```
fpos_t pos;  
FILE *fp = fopen("bit", "r");  
...  
fgetpos(fp, &pos);  
...  
fsetpos(fp, &pos);
```

fileno()

❑ synopsis

- `#include <stdio.h>`
- `int fileno(FILE *fp);`

❑ description

- examines the argument fp and returns the file descriptor associated with fp.

Temporary Files

■ synopsis

- `#include <stdio.h>`
- `char *tmpnam(char *s);`
- `FILE *tmpfile (void);`
- `char *tempnam(const char *dir, const char *pfx);`
- `#include <stdlib.h>`
- `char *mktemp(char *template);`
- `int mkstemp(char *template);`

Temporary Files

■ description

- `char *tmpnam(char *s);`
 - generates a unique temporary filename using the path prefix `P_tmpdir` defined in `<stdio.h>`
 - If the argument `s` is `NULL`, `tmpnam()` returns the address of an internal static area which holds the filename, which is overwritten by subsequent calls to `tmpnam()`
 - If `s` is not `NULL`, the filename is returned in `s`
 - return value
 - a pointer to the unique temporary filename
 - `NULL` if a unique name cannot be generated

Temporary Files

- FILE *tmpfile(void);
 - generates a unique temporary filename using the path prefix `P_tmpdir` defined in `<stdio.h>`
 - the temporary file is then opened in binary read/write (`w+b`) mode
 - the file will be automatically deleted when it is closed or the program terminates
 - return value
 - returns a stream descriptor
 - `NULL` if a unique filename cannot be generated or the unique file cannot be opened

Temporary Files

- `char *tempnam(const char *dir, const char *pfx)`
 - generates a unique temporary filename using up to five characters of pfx, if it is not NULL
 - The directory to place the file is searched for in the following order
 - The directory specified by the environment variable TMPDIR, if it is writable.
 - The directory specified by the argument dir, if it is not NULL.
 - The directory specified by P_tmpdir.
 - The directory /tmp.

Temporary Files

- The storage for the filename is allocated by malloc(), and so can be free'd by the function free().
- return value
 - a pointer to the unique temporary filename
 - NULL if a unique filename cannot be generated.
- `char *mktemp(char *template);`
 - generates a unique temporary filename from template.
 - the last six characters of template must be “XXXXXX” and these are replaced with a string that makes the filename unique.

Temporary Files

- template must not be a constant string.
- returns NULL on error, and template otherwise.
- int mkstemp(char *template);
 - generates a unique temporary filename from template like mktemp().
 - the file is then created with mode read/write.
 - returns the file descriptor of the temporary file or -1 on error

Variable Argument List

■ synopsis

- `#include <stdarg.h>`
- `va_list arglist;`
- `void va_start(va_list ap, last);`
- `type va_arg(va_list ap, type);`
- `void va_end(va_list ap);`

Variable Argument List

■ description

- `va_start()`

- initializes `ap` for subsequent use by `va_arg` and `va_end`, and must be called first
- The parameter `last` is the name of the last parameter before the variable argument list, i.e., the last parameter of which the calling function knows the type

- `va_arg()`

- an expression that has the type and value of the next argument in the call

Variable Argument List

- Each call to va_arg modifies ap so that the next call returns the next argument
 - The first use of the va_arg macro after that of the va_start macro returns the argument after last
 - Successive invocations return the values of the remaining arguments.
- va_end()
 - handles a normal return from the function whose variable argument list was initialized by va_start

Variable Argument List

Example

```
#include <stdio.h>
#include <stdarg.h>

int sum(int val, ...)
{
    va_list arglist;
    int arg, total;
    total = val;
    va_start(arglist, val);
    while ((arg = va_arg(arglist, int)) != 0)
        total += arg;
```

Variable Argument List

```
va_end(arglist);
return (total);
}

main()
{
    printf("Total sum: %d\n", sum(1,2,3,4,5,0));
}
```

Example – copy

```
#include <stdio.h>
#define MAXBUF 1024

main(int argc, char *argv[])
{
    FILE *source, *dest;
    char buf[MAXBUF];
    int count;
    if (argc != 3) {
        fprintf(stderr, "Usage: %s source destination\n",
                argv[0]);
        exit(1);
    }
```

Example - copy

```
if ((source = fopen(argv[1], "r")) == NULL) {  
    fprintf(stderr, "Can't open %s : ", argv[1]);  
    perror("");  
    exit(1);  
}  
if ((dest = fopen(argv[2], "w")) == NULL) {  
    fprintf(stderr, "Can't open %s : ", argv[2]);  
    perror("");  
    exit(1);  
}
```

Example - copy

```
while ((count = fread(buf, 1, MAXBUF, source))) {  
    if (fwrite(buf, 1, count, dest) == 0) {  
        perror("fwrite");  
        exit(1);  
    }  
}  
if (ferror(source)) {  
    perror("fread");  
    exit(1);  
}  
fclose(source);  
fclose(dest);  
}
```