

# Environment of a Process

---

# main Function

---

## ■ main()

- `int main(int argc, char *argv[]);`
  - `argc` : the number of command-line arguments
  - `argv` : an array of pointers to the arguments
- when a C program is started by the kernel, a special start-up routine is called before the `main()` is called.
- the start-up routine takes values from the kernel (the command-line arguments and the environment) and sets things up.

---

# Process Termination

---

## ❑ normal termination

- return from main
- calling exit
- calling \_exit

## ❑ abnormal termination

- calling abort
- terminated by a signal

---

# Process Termination

---

## exit()

- `#include <stdlib.h>`
- `void exit(int status);`
  - causes normal program termination and the value of status is returned to the parent
  - All functions registered with `atexit()` and `on_exit()` are called in the reverse order of their registration
  - all open streams are flushed and closed.

---

# Process Termination

---

## ■ atexit()

- `#include <stdlib.h>`
- `int atexit(void (*function)(void));`
  - registers the given function to be called at normal program termination
  - no arguments are passed
  - up to 32 functions can be registered
  - returns the value 0 if successful; otherwise the value -1 is returned

---

# Process Termination

---

## ■ on\_exit()

- `#include <stdlib.h>`
- `int on_exit(void (*function)(int , void *), void *arg);`
  - registers the given function to be called at normal program termination, whether via `exit()` or via return from the program's main.
  - The function is passed the argument to `exit()` and the `arg` argument from `on_exit()`
  - The integer argument of function is the exit status.

---

# Process Termination

---

## ▀ \_exit()

- #include <unistd.h>
- void \_exit(int status);
  - terminates the calling process immediately
  - Any open file descriptors belonging to the process are closed
  - status is returned to the parent process as the process's exit status

---

# Process Termination

---


```
#include <stdio.h>
#include <stdlib.h>

void exit1(void)
{
    puts("exit1");
}

void exit2(void)
{
    puts("exit2");
}

main()
{
    atexit(exit1);
    atexit(exit2);
    exit(3);
}
```

# Start and Termination of a C program



---


# Environment List

---


## ■ environment list

- `extern char **environ;`
- `int main(int argc, char *argv[], char *envp[])`
- each environment consists of “name=value”
- used to go through the entire environment.


# Environment List



# Memory Layout of a C Program



# Process Image



# Process Image

## ■ TEXT

- user가 작성한 instruction들이 들어 가는 영역이다.
- text는 항상 같은 virtual address에 load된다.
- text는 공유될 수도 있고 안 될 수도 있지만 default는 공유된다

## ■ DATA

- static variable과 global variable을 저장한다.
- DATA segment 부분은 initialized global variable을 가지는 initialized data 부분과 uninitialized global variable을 저장하는 uninitialized data 부분으로 나뉜다.

## ■ STACK

- stack 부분은 local variable, function의 argument, function의 return value를 저장한다
- data의 끝부분과 stack의 제일 윗부분 사이를 reference하면 memory fault 가 발생한다.

## ■ HEAP

- malloc(), calloc() 등의 인터페이스로 동적 할당되는 부분

## ■ USER AREA

- OS의 주소공간
- process가 수행 중일 때만 필요한 정보를 가지고 있다.
- open file list

# Process Image

```
#include <stdio.h>
#include <stdlib.h>
extern etext,edata,end;

char a='a';
int b=1;
int e;int f;
int func(int arg)
{
    int f=10;
    printf("addr of arg : %p:%d\n",
           &arg,arg);
    printf("addr of f : %p :%d\n",&f,f);
}

main()
{
    static int c=10;
    static int d;
    int *dynamic;
    printf("end of text %p\n", &etext);
    printf("addr of a : %p :%c\n",&a,a);
    printf("addr of b : %p :%d\n",&b,b);
    printf("addr of c : %p :%d\n",&c,c);
    printf("end of initialized data %p\n", &edata);
    printf("addr of d : %p :%d\n",&d,d);
    printf("addr of e : %p :%d\n",&e,e);
    printf("addr of f : %p :%d\n",&f,f);
    printf("end of uninitialized data %p\n", &end);
    func(5);
    dynamic=(int *)malloc(8);
    *dynamic=100;
    printf("addr of dynamic : %p :%d\n",dynamic,*dynamic);
}
```

---

# Shared Libraries

---

## ❑ shared libraries

- remove the common library routine somewhere in memory that all processes reference
- the executable file does not contain the library object code, but only a reference to the library name.
- when the program is loaded, program interpreter (ld.so)
  - analyzes the library names in the executable file
  - locates the library in the system's directory tree
  - makes the requested code available to the executing process.

---

# Memory Allocation

---

## ■ library functions

- `#include <stdlib.h>`
- `void *calloc(size_t nmemb, size_t size);`
  - allocates memory for an array of nmemb elements of size bytes each
  - returns a pointer to the allocated memory
  - The memory is set to zero
- `void *malloc(size_t size);`
  - allocates size bytes and returns a pointer to the allocated memory
  - The memory is not cleared

---

# Memory Allocation

---

- `void *realloc(void *ptr, size_t size);`
  - changes the size of the memory block pointed to by ptr to size bytes
  - The contents will be unchanged to the minimum of the old and new sizes
  - newly allocated memory will be uninitialized
  - If ptr is NULL, the call is equivalent to `malloc(size)`
  - if size is equal to zero, the call is equivalent to `free(ptr)`
- `void free(void *ptr);`
  - frees the memory space pointed to by ptr

---

# Memory Allocation

---

- `void *alloca( size_t size);`
  - allocates `size` bytes of space in the stack frame of the caller
  - this temporary space is automatically freed on return
- return value
  - returns a pointer to the allocated memory on success
  - the returned pointer of realloc may be different from `ptr` which is passed to realloc as an argument
  - returns `NULL` if the request fails

---

# Memory Allocation

---

## ■ system calls

- `#include <unistd.h>`
- `int brk(void *end_data_segment);`
  - sets the end of the data segment to the value specified by `end_data_segment`
  - `end_data_segment` must be greater than end of the text segment and it must be 16kB before the end of the stack
  - On success, `brk` returns zero
  - On error, -1 is returned

---

# Memory Allocation

---

- `void *sbrk(ptrdiff_t increment);`
  - C library wrapper(actually, not a system call)
  - increments the program's data space by increment bytes
  - returns a pointer to the start of the new area
  - On error, -1 is returned
- `brk` and `sbrk` should not be used with `malloc`, `calloc`, `realloc`

---

# Memory Allocation

---

## example

```
main()
{
    char *p;
    printf("before malloc: %p\n", sbrk(0));
    p = malloc(4096);
    printf("after malloc: %p\n", sbrk(0));
    free(p);
    printf("after free: %p\n", sbrk(0));
}
```

---

# Environment Variables

---

## ❑ getenv() and putenv()

- `#include <stdlib.h>`
- `char *getenv(const char *name);`
  - searches the environment list for a string that matches the string pointed to by name
  - returns a pointer to the value in the environment, or NULL if there is no match

---

# Environment Variables

---

- `int putenv(const char *string);`
  - adds or changes the value of environment variables
  - The argument string is of the form name=value
  - If name does not already exist in the environment, then string is added to the environment
  - If name does exist, then the value of name in the environment is changed to value
  - returns zero on success, or -1 if an error occurs

---

# Environment Variables

---

## ■ setenv() and unsetenv()

- `#include <stdlib.h>`
- `int setenv(const char *name, const char *value, int overwrite);`
  - adds the variable name to the environment with the value value, if name does not already exist
  - If name does exist in the environment, then its value is changed to value if overwrite is non-zero; if overwrite is zero, then the value of name is not changed.
- `void unsetenv(const char *name);`
  - deletes the variable name from the environment

---

# Environment Variables

---

## example

```
#include <stdio.h>
#include <stdlib.h>

main()
{
    printf("Home directory is %s\n",
           getenv("HOME"));
    putenv("HOME=/");
    printf("New home directory is %s\n",
           getenv("HOME"));
}
```

---

# setjmp and longjmp

---

## ■ setjmp( ) and longjmp( )

- `#include <setjmp.h>`
- `int setjmp(jmp_buf env);`
  - useful for dealing with errors and interrupts encountered in a low-level subroutine of a program
  - saves the stack context/environment in env for later use by longjmp()
  - The stack context will be invalidated if the function which called setjmp() returns
  - return 0 if returning directly, and non-zero when returning from longjmp() using the saved context

---

# setjmp and longjmp

---

- `void longjmp(jmp_buf env, int val);`
  - restores the environment saved by the last call of `setjmp()` with the corresponding `env` argument
  - After `longjmp()` is completed, program execution continues as if the corresponding call of `setjmp()` had just returned the value `val`
  - cannot cause 0 to be returned.
  - If `longjmp` is invoked with a second argument of 0, 1 will be returned instead.

# setjmp and longjmp

```
#include <setjmp.h>
#include <stdio.h>
#include <stdlib.h>
static void          f1(int, int, int);
static void          f2(void);
static jmp_buf       jmpbuffer;
int main(void)
{
    int      count;
    register int val;
    volatile int sum;
    count = 2; val = 3; sum = 4;
    if (setjmp(jmpbuffer) != 0) {
        printf("after longjmp:\n");
        printf("count= %d, val=%d, sum=%d\n",
               count, val, sum);
        exit(0);
    }
    count = 97; val = 98; sum = 99;
    /* changed after setjmp, before longjmp */
    /*
     * f1(count, val, sum); /* never returns
     */
}
static void f1(int i, int j, int k)
{
    printf("in f1(): count = %d, val = %d,
           sum = %d\n", i, j, k);
    f2();
}
static void f2(void)
{
    longjmp(jmpbuffer, 1);
}
```

\$ gcc testjmp.c

\$ gcc -O testjmp.c

---

# Resource Limit and Usage

---

## ❑ getrlimit() and setrlimit()

- #include <sys/time.h>
- #include <sys/resource.h>
- #include <unistd.h>
  
- int getrlimit (int resource, struct rlimit \*rlim);
- int setrlimit (int resource, const struct rlimit \*rlim);

---

# Resource Limit and Usage

---

- resource
  - RLIMIT\_CPU /\* CPU time in seconds (SIGXCPU)\*/
  - RLIMIT\_FSIZE /\* Maximum filesize (SIGXFSZ)\*/
  - RLIMIT\_DATA /\* max data size \*/
  - RLIMIT\_STACK /\* max stack size \*/
  - RLIMIT\_CORE /\* max core file size \*/
  - RLIMIT\_RSS /\* max resident set size \*/
    - if physical memory is tight, the kernel takes memory from processes that exceed their RSS

---

# Resource Limit and Usage

---

- RLIMIT\_NPROC /\* max number of processes per real user ID\*/
- RLIMIT\_NOFILE /\* max number of open files per process\*/

- rlim

```
struct rlimit
{
    int rlim_cur; /* soft limit: current limit */
    int rlim_max; /* hard limit: maximum value for rlim_cur */
};
```

- RLIM\_INFINITY : resource is unlimited

---

# Resource Limit and Usage

---

## ▣ getrusage()

- #include <sys/time.h>
- #include <sys/resource.h>
- #include <unistd.h>
  
- int getrusage (int who, struct rusage \*usage);
  - who : RUSAGE\_SELF, RUSAGE\_CHILDREN

---

# Resource Limit and Usage

---

```
struct rusage
{
    struct timeval ru_utime; /* user time used */
    struct timeval ru_stime; /* system time used */
    long ru_maxrss; /* maximum resident set size */
    long ru_ixrss; /* integral shared memory size */
    long ru_idrss; /* integral unshared data size */
    long ru_isrss; /* integral unshared stack size */
    long ru_minflt; /* page reclaims */
    long ru_majflt; /* page faults */
    long ru_nswap; /* swaps */
```

---

# Resource Limit and Usage

---

```
long ru_inblock;      /* block input operations */
long ru_oublock;     /* block output operations */
long ru_msgsnd;      /* messages sent */
long ru_msgrcv;      /* messages received */
long ru_nsignals;    /* signals received */
long ru_nvcsw;       /* voluntary context switches */
long ru_nivcsw;      /* involuntary context switches */
};
```