

# Interprocess Communication



Pipes



FIFOs



System V IPC

- Message Queues
- Shared Memory
- Semaphores

# pipe

## IPC using regular files

- unrelated processes can share
- fixed size
- lack of synchronization

## IPC using pipes

- for transmitting data between related processes
- can transmit an unlimited amount of data
- automatic synchronization on open()

# Pipe in a Unix Shell

💡 In a UNIX shell, the pipe symbol is: | (the vertical bar)

💡 In a shell, UNIX pipes look like:

```
$ ls -alg | more
```

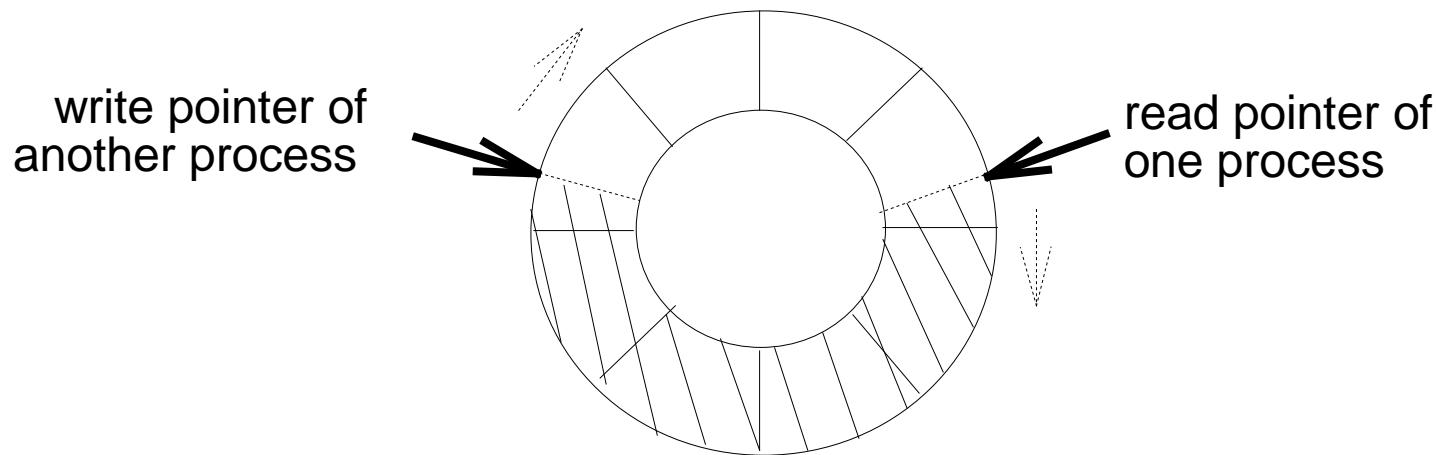
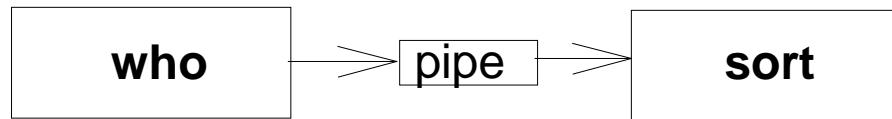
- where the standard output of the program at the left (i.e., the producer) becomes the standard input of the program at the right (i.e., the consumer).

💡 We can have longer pipes:

```
$ ps -ef | sort | more
```

# Example

```
$ who | sort
```



# pipe

## Data transmitting

- data is written into pipes using the write() system call
- data is read from a pipe using the read() system call
- automatic blocking when full or empty

## Types of pipes

- (unnamed) pipes
- named pipes (FIFOs)

# pipe(1/4)

● In UNIX, pipes are the oldest form of IPC.

● Limitations of Pipes:

- Half duplex (data flows in one direction)
- Can only be used between processes that have a common ancestor  
(Usually used between the parent and child processes)
- Processes cannot pass pipes and must inherit them from their parent
- If a process creates a pipe, all its children will inherit it

# pipe(2/4)

```
#include <unistd.h>

int pipe(int fd[2])
```

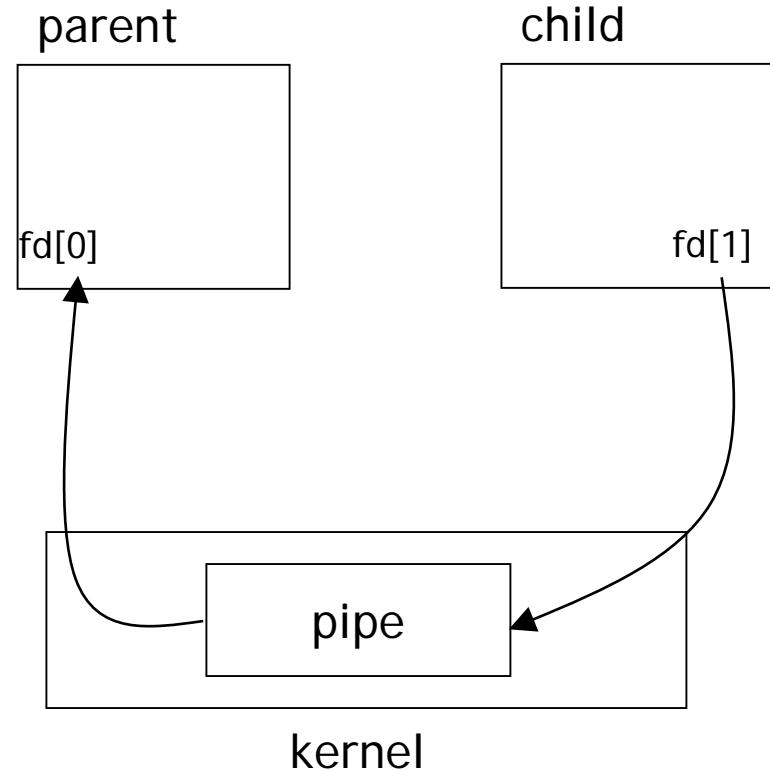
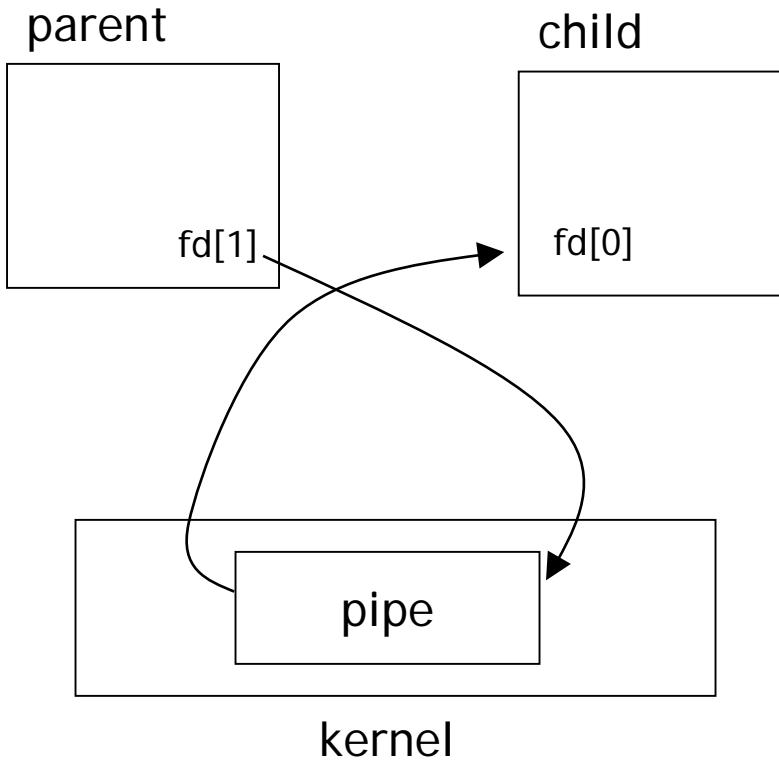
**Returns:** 0 if OK, -1 on error

- ➊ Two file descriptors are returned through the *fd* argument
  - *fd[0]*: can be used to read from the pipe, and
  - *fd[1]*: can be used to write to the pipe
- ➋ Anything that is written on *fd[1]* may be read by *fd[0]*.
  - This is of no use in a single process.
  - However, between processes, it gives a method of communication
- ➌ The pipe() system call gives parent-child processes a way to communicate with each other.

# pipe(3/4)

parent → child:  
parent closes fd[0]  
child closes fd[1]

parent ← child:  
parent closes fd[1]  
child closes fd[0]



# pipe(3/4)

 Read from a pipe with write end closed: (fd[1]가 close된 경우)

- returns 0 to indicate EOF

 Write to a pipe with read end closed: (fd[0]가 close된 경우)

- SIGPIPE generated,
- write() returns error (errno == EPIPE)

```
#include <stdio.h> // pipe.c

#define READ  0
#define WRITE 1

char* phrase = "Stuff this in your pipe and smoke it";

int main( ) {
    int fd[2], bytesRead;
    char message[100];
    pipe(fd);
    if (fork() == 0) { // child
        close(fd[READ]);
        write(fd[WRITE], phrase, strlen(phrase)+1);
        fprintf(stdout, "[%d, child] write completed.\n", getpid());
        close(fd[WRITE]);
    }
    else { // parent
        close(fd[WRITE]);
        bytesRead = read(fd[READ], message, 100);
        fprintf(stdout, "[%d, parent] read completed.\n", getpid());
        printf("Read %d bytes: %s\n", bytesRead,message);
        close(fd[READ]);
    }
}
```

# System V IPC

## Message Queues

- Send and receive amount of data called “messages”.
- The sender classifies each message with a type.

## Shared Memory

- Shared memory allows two or more processes to share a given region of memory.
- Readers and writers may use semaphore for synchronization.

## Semaphores

- Process synchronization and resource management
- For example, a semaphore might be used to control access to a device like printer.

# Identifier & Key

- ❶ Identifier : each IPC structure has a nonnegative integer
- ❷ Key: when creating an IPC structure, a *key* must be specified (**key\_t**)  
`id = xxxget(key, ...)`
- ❸ How to access the same IPC? → key in a common header
  - Define a key in a common header
  - Client and server agree to use that key
  - Server creates a new IPC structure using that key
  - Problem when the key is already in use
    - (msgget, semget, shmget returns error)
    - Solution: delete existing key, create a new one again!

# IPC System Calls

## **msg/sem/shm get**

- Create new or open existing IPC structure.
- Returns an IPC identifier

## **msg/sem/shm ctl**

- Determine status, set options and/or permissions
- Remove an IPC identifier

## **msg/sem/shm op**

- Operate on an IPC identifier
- For example(Message queue)
  - add new msg to a queue (msgsnd)
  - receive msg from a queue (msgrcv)

# Permission structure

- ipc\_perm is associated with each IPC structure.
- Defines the permissions and owner.

```
struct ipc_perm {  
    uid_t uid;      /* owner's effective user id */  
    gid_t gid;      /* owner's effective group id */  
    uid_t cuid;     /* creator's effective user id */  
    gid_t cgid;     /* creator's effective group id */  
    mode_t mode;    /* access modes */  
    ulong seq;      /* slot usage sequence number */  
    key_t key;      /* key */  
};
```

# message queue



## Linked list of messages

- Stored in kernel
- Identified by message queue identifier (in kernel)



## msgget

- Create a new queue or open existing queue.



## msgsnd

- Add a new message to a queue



## msgrcv

- Receive a message from a queue
- Fetching order: based on type

# message queue

- Each queue has a structure

```
struct msqid_ds {  
    struct ipc_perm msg_perm;  
    struct msg *msg_first; /* ptr to first msg on queue */  
    struct msg *msg_last; /* ptr to last msg on queue */  
    ulong msg_cbytes; /* current # bytes on queue */  
    ulong msg_qnum; /* # msgs on queue */  
    ulong msg_qbytes; /* max # bytes on queue */  
    pid_t msg_lspid; /* pid of last msgsnd() */  
    pid_t msg_lrpid; /* pid of last msgrcv() */  
    time_t msg_stime; /* last-msgsnd() time */  
    time_t msg_rtime; /* last-msgrcv() time */  
    time_t msg_ctime; /* last-change time */  
};
```

- We can get the structure using msgctl() function.
- Actually, however, we don't need to know the structure in detail. 17

# msgget()

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>

int msgget(key_t key, int flag);
    Returns: msg queue ID if OK, -1 on error
```

- >Create new or open existing queue
- flag : ipc\_perm.mode
- Example

```
msg_qid = msgget(DEFINED_KEY, IPC_CREAT | 0666);
```

# msgctl()

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>

int msgctl(int msqid, int cmd, struct msqid_ds *buf);
    Returns: 0 if OK, -1 on error
```

- ④ Performs various operations on a queue
- ④ **cmd = IPC\_STAT:**  
fetch the msqid\_ds structure for this queue, storing it in buf
- ④ **cmd = IPC\_SET:**  
set the following four fields from buf: msg\_perm.uid, msg\_perm.gid, msg\_perm.mode, and msg\_qbytes
- ④ **cmd = IPC\_RMID:**  
remove the message queue.

# msgsnd()

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>

int msgsnd(int msqid, const void *ptr, size_t nbytes, int flag);
    Returns: 0 if OK, -1 on error
```

- msgsnd() places a message at the end of the queue.

- ptr: pointer that points to a message
- nbytes: length of message data
- if flag = IPC\_NOWAIT: IPC\_NOWAIT is similar to the nonblocking I/O flag for file I/O.

- Structure of messages

```
struct mymsg {
    long mtype;          /* positive message type */
    char mtext[512];     /* message data, of length nbytes */
};
```

# msgrecv()

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>

int msgrecv(int msqid, void *ptr, size_t nbytes, long type, int flag);
    Returns: data size in message if OK, -1 on error
```

- msrecv() retrieves a message from a queue.
- type == 0: the first message on the queue is returned
- type > 0: the first message on the queue whose message type equals type is returned
- type < 0: the first message on the queue whose message type is the lowest value less than or equal to the absolute value of type is returned
- flag may be given by IPC\_NOWAIT

# sender.c

```
#include <stdio.h> // sender.c
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>

#define DEFINED_KEY 0x10101010

main(int argc, char **argv)
{
    int msg_qid;
    struct {
        long mtype;
        char content[256];
    } msg;

    fprintf(stdout, "=====SENDER=====\\n" );

    if((msg_qid = msgget(DEFINED_KEY, IPC_CREAT | 0666)) < 0) {
        perror("msgget: ");
        exit(-1);
    }

    msg.mtype = 1;
    while(1) {
        memset(msg.content, 0x0, 256);
        gets(msg.content);
        if(msgsnd(msg_qid, &msg, sizeof(msg.content), 0) < 0) {
            perror("msgsnd: ");
            exit(-1);
        }
    }
}
```

# receiver.c

```
#include <stdio.h> // receiver.c
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/msg.h>

#define DEFINED_KEY 0x10101010

main(int argc, char **argv)
{
    int msg_qid;
    struct {
        long mtype;
        char content[256];
    } msg;

    fprintf(stdout, "=====RECEIVER=====\\n" );
    if((msg_qid = msgget(DEFINED_KEY, IPC_CREAT | 0666)) < 0) {
        perror("msgget: ");
        exit(-1);
    }

    while(1) {
        memset(msg.content, 0x0, 256);
        if(msgrcv(msg_qid, &msg, 256, 0, 0) < 0) {
            perror("msgrcv: ");
            exit(-1);
        }
        puts(msg.content);
    }
}
```

# Shared memory

- ➊ Allows multiple processes to share a region of memory
  - Fastest form of IPC: no need of data copying between client & server
  
- ➋ If a shared memory segment is attached
  - It become a part of a process data space, and shared among multiple processes
  
- ➌ Readers and writers may use semaphore to
  - synchronize access to a shared memory segment

# Shared memory structure

- Each shared memory has a structure

```
struct shmid_ds {  
    struct ipc_perm shm_perm;  
    struct anon_map *shm_amp; /* pointer in kernel */  
    int shm_segsz;           /* size of segment in bytes */  
    ushort shm_lkcnt;       /* # of times segment is being locked */  
    pid_t shm_lpid;         /* pid of last shmop() */  
    pid_t shm_cpid;         /* pid of creator */  
    ulong shm_nattch;       /* # of current attaches */  
    ulong shm_cnattch;      /* used only for shminfo() */  
    time_t shm_atime;       /* last-attach time */  
    time_t shm_dtime;       /* last-detach time */  
    time_t shm_ctime;       /* last-change time */  
};
```

- We can get the structure using `shmctl()` function.
- Actually, however, we don't need to know the structure in detail.

# shmget()

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>

int shmget(key_t key, int size, int flag);
    Returns: shared memory ID if OK, -1 on error
```

- ➊ Obtain a shared memory identifier
- ➋ size: is the size of the shared memory segment
- ➌ flag: ipc\_perm.mode
- ➍ Example

```
shmid = shmget(key, size, PERM|IPC_CREAT|IPC_EXCL|0666);
```

# shmctl

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>

int shmctl(int shmid, int cmd, struct shmid_ds *buf);
    Returns: 0 if OK, -1 on error
```

- ④ Performs various shared memory operations
- ④ **cmd = IPC\_STAT:**  
fetch the shmid\_ds structure into *buf*
- ④ **cmd = IPC\_SET:**  
set the following three fields from *buf*: *shm\_perm.uid*, *shm\_perm.gid*, and *shm\_perm.mode*
- ④ **cmd = IPC\_RMID:**  
remove the shared memory segment set from the system

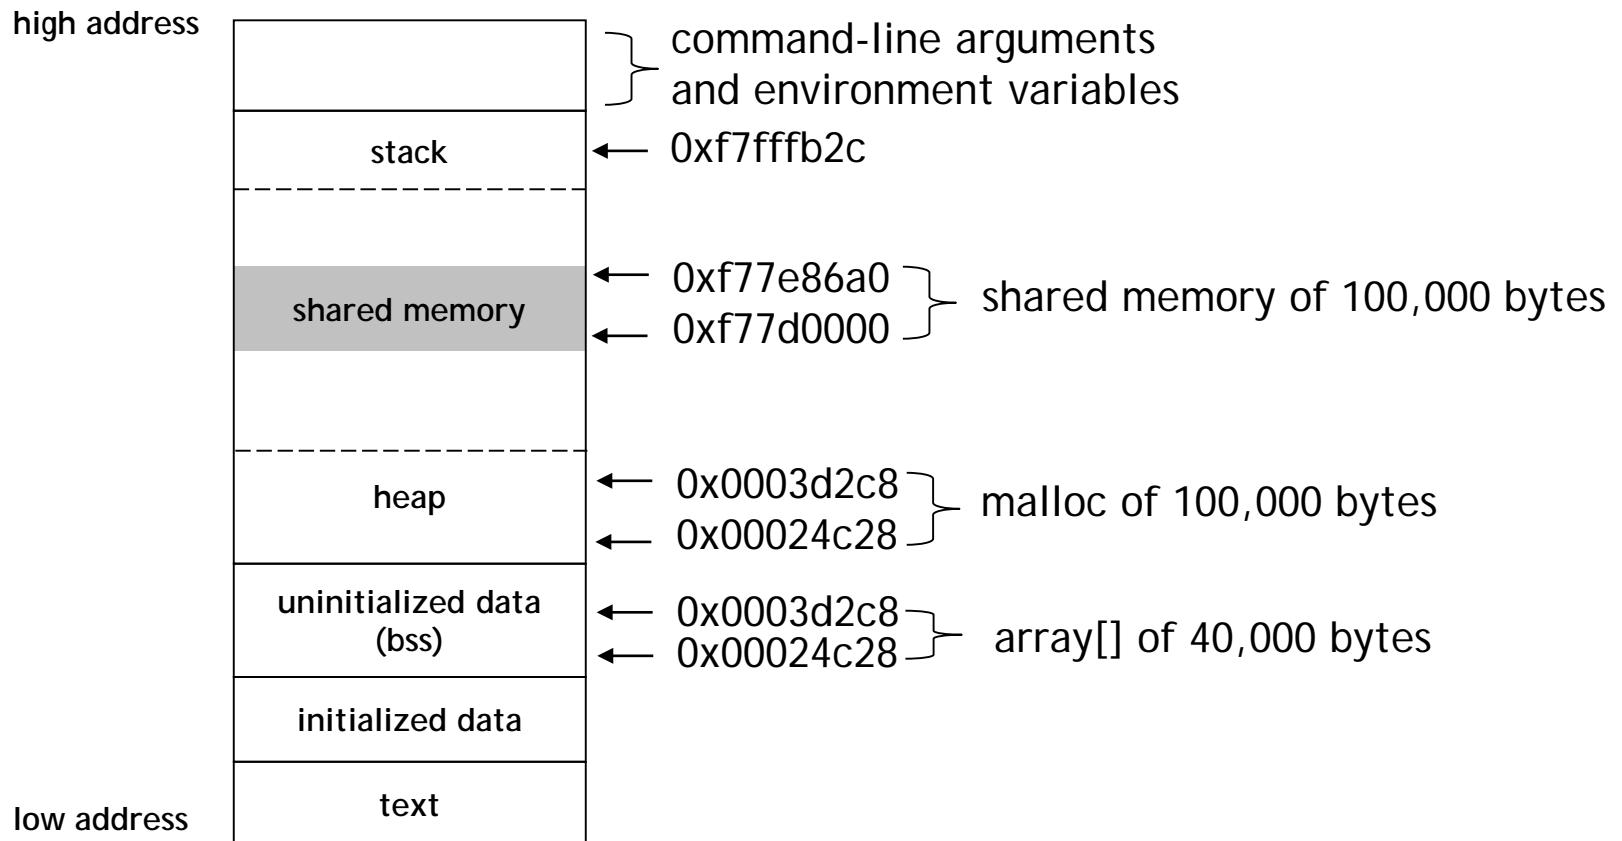
# shmat()

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>

void *shmat (int shmid, void *addr, int flag);
    Returns: pointer to shared memory segment if OK, -1 on error
```

- Attached a shared memory to an address
- flag = SHM\_RDONLY: the segment is read-only
- addr==0: at the first address selected by the kernel (recommended!)
- addr!=0: at the address given by addr

# Memory Layout



# shmdt()

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>

void shmdt (void *addr);
```

Returns: 0 if OK, -1 on error

- Detach a shared memory segment

```
#include <sys/types.h> // shm.c
#include <sys/ipc.h>
#include <sys/shm.h>

#define ARRAY_SIZE    100000
#define MALLOC_SIZE   100000
#define SHM_SIZE      100000

err_sys(char *p) { perror(p); exit(-1); }

char array[ARRAY_SIZE]; /* uninitialized data = bss */

int main(void) {
    int shmid; char *ptr, *shmptr;

    printf("array[] from %x to %x\n", &array[0], &array[ARRAY_SIZE]);
    printf("stack around %x\n", &shmid);

    if ((ptr = malloc(MALLOC_SIZE)) == NULL) err_sys("malloc error");
    printf("malloced from %x to %x\n", ptr, ptr+MALLOC_SIZE);

    if ((shmid = shmget(0x01010101, SHM_SIZE, IPC_CREAT | 0666)) < 0)
        err_sys("shmget error");

    if ((shmptr = shmat(shmid, 0, 0)) == (void *) -1) err_sys("shmat error");
    printf("shared memory attached from %x to %x\n", shmptr, shmptr+SHM_SIZE);

    // if (shmctl(shmid, IPC_RMID, 0) < 0) err_sys("shmctl error");

    exit(0);
}
```

# shm.c

## 실행 결과

```
$ shm
```

```
array[] from 20bd8 to 39278
```

```
stack around ff9fffa9c
```

```
mallocoed from 39288 to 51928
```

```
shared memory attached from ff260000 to ff2786a0
```

```
$ ipcs -ma
```

```
IPC status from <running system> as of 2009년 11월 14일 토요일 오후 06시 45분 09초
```

T	ID	KEY	MODE	OWNER	GROUP	CREATOR	CGROUP	NATTCH	SEGSZ	CPID
LPID	ATIME	DTIME	CTIME							

```
Shared Memory:
```

m	1	0x1010101	--rw-rw-rw-	yshin	prof	yshin	prof	0	100000	2413 2465
		18:42:50	18:42:50	18:32:37						

```
$ ipcrm -m 1
```

```
$ ipcs -ma
```

```
#include <stdio.h>                      // shm1.c
#include <stdlib.h>
#include <string.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
#define KEY 0x101010
#define SIZE 1024

main()
{
    int shmid;
    char *shmaddr;

    if ((shmid=shmget(KEY, SIZE, IPC_CREAT|0666)) == -1) {
        perror("shmid failed");
        exit(1);
    }

    if ((shmaddr=shmat(shmid, NULL, 0)) == (void *)-1) {
        perror("shmat failed");
        exit(1);
    }
    strcpy((char *)shmaddr, "HELLO KIM!!");

    if (shmdt(shmaddr) == -1) {
        perror("shmdt failed");
        exit(1);
    }
}
```

```
#include <stdio.h>           // shm2.c
#include <stdlib.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>

#define KEY 0x101010
#define SIZE 1024

main()
{
    int shmid;
    void *shmaddr;

    if((shmid=shmget(KEY, SIZE, IPC_CREAT|0666)) == -1) {
        perror("shmid failed");
        exit(1);
    }

    if((shmaddr=shmat(shmid, NULL, 0)) == (void *)-1) {
        perror("shmat failed");
        exit(1);
    }
    printf("data read from shared memory : %s\n", (char
*)shmaddr);

    if(shmdt(shmaddr) == -1) {
        perror("shmdt failed");
        exit(1);
    }

    if(shmctl(shmid, IPC_RMID, 0) == -1) {
        perror("shmctl failed");
        exit(1);
    }
}
```

# semaphore

- ❶ A counter to provide access to shared data object for multiple processes (복수의 프로세스가 데이터를 공유하는데 사용하는 카운터)
- ❷ To obtain a shared resource:
  - 1. Test semaphore that controls the resource (확인하여)
  - 2. If value > 0, value--, grant use (양수이면, 감소시키고 사용하고)
  - 3. If value == 0, sleep until value > 0 (0이면 기다림)
  - 4. Release resource, value ++ (다 쓴 후에는 다시 양수로 만듦)
- ❸ Step 1, 2 must be an **atomic operation**

# Semaphore structure

- Each semaphore has a structure

```
struct semid_ds {  
    struct ipc_perm sem_perm;  
    struct sem *sem_base; /*ptr to first semaphore in set */  
    ushort sem_nsems; /* # of semaphors in set */  
    time_t sem_otime; /* last-semop() time */  
    time_t sem_ctime; /* last-change time */  
};  
  
struct sem {  
    ushort semval; /* semaphore value, always >= 0 */  
    pid_t sempid; /* pid for last operation */  
    ushort semncnt; /* # processes awaiting semval > currval */  
    ushort semzcnt; /* # processes awaiting semval = 0 */  
};
```

- We can get the structure using `semctl()` function.
- Actually, however, we don't need to know the structure in detail. 36

# semget()

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>

int semget(key_t key, int nsems, int flag);
    Returns: semaphore ID if OK, -1 on error
```

- Obtain a semaphore ID
- nsems: sem\_nsens (# of semaphores in set)
- flag: ipc\_perm.mode

# semctl()

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>

int semctl(int semid, int semnum, int cmd, union semun arg);

union semun {
    int             val;    /* for SETVAL */
    struct semid_ds *buf;  /* for IPC_START and IPC_SET */
    ushort         *array; /* for GETALL and SETALL */
};
```

- To use semaphore, please refer to the textbook and manuals related semaphore.

# ipcs, ipcrm

## ④ ipcs: System V IPC의 상태를 확인하는 명령어

- `$ ipcs` // IPC 정보를 확인 (q, m, s 모두)
- `$ ipcs -q ($ ipcs -qa)` // Message Queue 정보를 확인
- `$ ipcs -m ($ ipcs -ma)` // Shared Memory 정보를 확인
- `$ ipcs -s ($ ipcs -sa)` // Semaphore 정보를 확인

## ⑤ ipcrm: 정의된(생성된) IPC를 삭제함

- `$ ipcrm -q id` // Message Queue를 삭제
- `$ ipcrm -m id` // Shared Memory를 삭제
- `$ ipcrm -s id` // Semaphore를 삭제